Qwen

Architecting the Cybernetic Hero: A Multi-Layered Framework for Measuring Ethical Impact Through Language, Evidence, and Identity

Ethical Foundations: The Non-Negotiable Hard Gates of Anti-Barbarism and Anti-Greed

The foundation of the Cybernetic Hero reputation framework is not a simple metric of effectiveness but a robust system of ethical governance designed to distinguish morally grounded heroism from mere tactical success. This is achieved through the implementation of two "hard gates": the anti-barbaric invariant and the anti-greed invariant . These invariants function as absolute filters that disqualify any action from consideration as heroic, irrespective of its outcome or perceived courage. They serve to codify a moral code into the very logic of the scoring algorithm, preventing the system from rewarding behavior that is effective yet fundamentally unethical. The first gate, the anti-barbaric invariant, establishes a clear boundary between heroic action and cruelty. It is triggered when an act crosses a predefined threshold of toxicity, encompassing behaviors such as torture, desecration, scapegoating, or the betrayal of the defenseless . This principle aligns with long-standing philosophical and legal distinctions between civility and barbarism, particularly within the domain of international humanitarian law, which seeks to define who the laws of war are meant to protect [11] [33]. Operationally, this gate relies on a computational toxicity metric, T, which must remain below a stringent hero-grade threshold, $T_{\text{hero_max}}$. While modern tools like Perspective API offer standardized methods for measuring rudeness and disrespect in language [24] [49], their limitations are well-documented. They may fail to capture nuanced forms of symbolic violence or context-specific cruelty, especially across different languages and cultures [25] [78]. Therefore, achieving a scientifically valid application of this gate requires more than off-the-shelf technology; it necessitates either highly sophisticated, context-aware models trained on extensive ethical corpora or, more pragmatically, a hybrid approach combining automated flagging with human-in-the-loop expert review to adjudicate borderline cases [45] [108].

The second hard gate, the anti-greed invariant, addresses the motivation behind an act. An action driven primarily by greed—defined as an acquisitive desire for excessive, extractive rewards—is disqualified from being considered heroic [19]. This mechanism is crucial for distinguishing heroes whose motivations are rooted in duty, care, or justice from "warlord archetypes" whose actions, however successful, are instruments of domination and exclusion . The challenge lies in quantifying "greed." Unlike toxicity, which can be partially measured linguistically, greed is a behavioral and motivational trait. The framework proposes identifying greed-driven acts by analyzing whether rewards wildly exceed the service rendered or if the reputation itself is weaponized for blackmail or gatekeeping . This distinction echoes academic work that draws a line between corruption driven by "need" versus "greed," suggesting that citizens are more likely to resist the former [18]. In practice, this would require auditing the entire transactional and relational context surrounding an act, not just its public narrative. For instance, accepting a modest reward for saving lives might be legitimate compensation, whereas demanding exorbitant payment or using a reputation score to block access to essential resources would trigger the anti-greed gate. The user's clarification that all verified heroic deeds contribute equally over time underscores the gravity of this gate; a single, egregious act of greed could potentially taint a long history of otherwise exemplary service, reflecting a strong philosophical stance on the irredeemable nature of certain betrayals of trust [32]. The enforcement of these hard gates is the most critical and operationally demanding aspect of the entire framework. It cannot be fully automated without risking profound misclassification errors. The system's integrity depends on creating a transparent, auditable process where acts flagged by the algorithms are subject to rigorous, independent verification before being definitively admitted or rejected from the pool of eligible heroic contributions.

Quantitative Evaluation: The Word-Math Quality and Social-Impact Vectors

Once an act has passed the non-negotiable ethical hard gates, it enters the quantitative evaluation phase, governed by two powerful vectors: the Word-Math quality score (F) and the Social-Impact vector (S). Together, they form a dual-axis assessment that evaluates both the *narrative form* and the *tangible consequence* of a heroic deed, ensuring that only acts that are both narratively coherent and truthfully reported, and that produce demonstrably positive social outcomes, are rewarded with reputation points. The Word-Math quality score, represented by the formula $F=(1-y)(1-z)(1-T)KE$, serves as

a proxy for the epistemic and ethical quality of the narrative describing the act . Each component is meticulously designed to filter for specific attributes of high-quality, trustworthy communication. The repetition contamination metric, y, penalizes redundant phrasing, encouraging concise and focused storytelling [43] [104]. Topic drift, z, measures the coherence of the narrative, ensuring the description remains tightly centered on the heroic event and its context rather than veering into irrelevant tangents [71]. As established by the hard gates, the toxicity metric, T, must be kept at a very low level, representing a core requirement for any qualifying act . The kindness or prosocial tone metric, K, directly attempts to quantify the moral framing of the action, rewarding narratives that are protective, respectful, and uplifting [40]. Finally, and most critically, the evidentiality metric, E, ensures that all claims are backed by verifiable proof, grounding the entire system in empirical reality [67] [91]. High values for K and E are mandatory for a hero-grade act, encoding the principle that true heroism involves restraint and is demonstrable . The multiplication of these factors creates a multiplicative penalty system; a weakness in one area can significantly reduce the overall score, reflecting the holistic nature of evaluating a story's credibility and moral character.

Complementing the linguistic analysis of the F score is the Social-Impact vector, $S=(S_{\text{antistigma}},S_{\text{nonexclusion}},S_{\text{peacekeeping}},S_{\text{eco}})$. This vector moves beyond the narrative to assess the real-world consequences of the act, prioritizing contributions that build a better society. The four axes are carefully chosen to reflect contemporary values and pressing global challenges. High scores in `antistigma` and `nonexclusion` reward efforts that reduce prejudice against marginalized groups and ensure equitable access to resources and opportunities [14] [15]. Peacekeeping emphasizes de-escalation and the proportional use of force, while the `eco` axis prioritizes ecological repair and rejects ecocidal trade-offs [64] [75]. Measuring these dimensions presents significant methodological challenges. There is currently no consensus in the research community on the definition or measurement of stigma, leading to fragmented and highly context-specific tools [13]. Similarly, quantifying peacekeeping or eco-civic impact requires sophisticated, domain-specific analyses that may involve tracking policy changes, analyzing shifts in public discourse, or monitoring environmental data [16]. Despite these challenges, the framework proposes compressing the vector S into a scalar, $\bar{S}$, which can then be multiplied by the quality score F to generate an "Output Factor," $O=F\cdot\bar{S}$. This composite score effectively functions as a filter, requiring acts to meet minimum thresholds for both narrative quality ($F\geq F_{\text{hero_min}}$) and social value ($O\geq O_{\text{hero_min}}$) to be counted . This dual-filter mechanism is central to the system's scientific validity, as it prevents the amplification of charismatic but unverified or socially harmful behavior. Only acts that are not only well-told but also genuinely beneficial can earn reputation, structurally reinforcing a culture of accountability and community benefit .

Metric/Vector	Component	Description & Measurement Challenge
Word-Math Score (F)	Repetition Contamination (y)	Measures lexical redundancy and verbosity. [43] [104]
	Topic Drift (z)	Measures narrative coherence and focus. Computationally challenging to define precisely. [71]
	Toxicity (T)	Measures rude/disrespectful content. Requires context-aware models beyond standard APIs. [24] [78]
	Kindness (K)	Measures prosocial, respectful, and protective language. Subject to cultural and contextual bias. [40]
	Evidentiality (E)	Measures the presence of verifiable evidence ("EvidenceBundle"). Foundational for scientific validity. [67] [91]
Social-Impact Vector (S)	Antistigma	Reduces negative attitudes and discrimination. No consensus on definition or measurement tools. [13] [14]
	Nonexclusion	Protects access to basic services and avoids gatekeeping. Requires policy and access data. [15]
	Peacekeeping	De-escalates conflict and uses force proportionally. Difficult to quantify without conflict data. [64]
	Eco-Civic	Prioritizes ecological repair and sustainable development. Requires environmental and policy data. [75]

Governance Architecture: DID-Bound, Non-Transferable Reputation and Neurorights

The governance architecture of the Cybernetic Hero framework is designed to ensure the reputation score is durable, trustworthy, and resistant to systemic corruption. This layer is built upon three core principles: identity binding via Decentralized Identifiers (DID), the creation of a non-transferable reputation ledger, and the enforcement of a neurorights envelope to protect individual agency and privacy. By anchoring every reputation record to a specific DID, the system creates a persistent, cryptographically verifiable link between an identity and its historical record of actions [7] [22]. This eliminates pseudonymity and allows for a true audit trail of an individual's contributions over time. The use of W3C standards for DIDs and Verifiable Credentials (VCs) provides a robust technical foundation, enabling entities within the ecosystem to issue, hold, and verify credentials without relying on a central intermediary [1] [87] [106]. The "EvidenceBundles" that validate heroic acts would be expressed as VCs, providing a secure and privacy-preserving way to attest to the occurrence and nature of an action .

This cryptographic grounding is fundamental to the score's scientific validity, transforming it from a speculative rating into a formal, auditable record.

A cornerstone of this governance model is the principle of non-transferability. Reputation tokens minted for heroic acts are explicitly designed to be non-transferable, meaning they cannot be bought, sold, or delegated to another entity . This brilliant design choice directly addresses the user's concern about preventing the emergence of "warlord archetypes"—individuals who accumulate power and influence not through personal merit but through extraction and control . By making reputation a personal and inseparable asset tied to an individual's DID, the system enforces a strict meritocracy. An agent's standing is determined solely by their own verified, rights-respecting actions. This concept resonates with emerging ideas in blockchain-based economies, where "Reputation Capital" is seen as a viable alternative to traditional collateral, and where reputation should be context-specific rather than a monolithic, fungible asset [22] [127]. The inability to consolidate reputation across multiple identities or sell it for other forms of capital prevents the hoarding of influence and ensures that those with high scores have earned them through direct engagement with the world. This contrasts sharply with many current online systems where influence can be purchased or artificially inflated, undermining the authenticity of any reputation metric.

Finally, the framework incorporates a neurorights envelope to safeguard against manipulative technologies and protect the sanctity of the individual's inner life. The score is explicitly prohibited from using raw neural signals or emotional states as inputs . All scoring is predicated on "expressed, consented actions with EvidenceBundles" . This is a critical ethical and privacy boundary. It prevents the system from becoming a tool for coercive neuromarketing or psychological manipulation, where an agent could be scored based on their involuntary emotional responses. The focus remains squarely on observable, verifiable behavior. This aligns with broader discussions in AI ethics regarding the management of human-AI interaction and the responsibilities of developers for the societal impact of their creations [35] [36]. The combination of DID-based identity, non-transferable assets, and a neurorights-enforced boundary creates a governance structure that is not only technically sound but also philosophically aligned with principles of accountability, fairness, and personal sovereignty. It builds a reputation system that is not merely a database of points, but a formal testament to a unique identity's consistent pattern of ethical and impactful action.

Aggregation Model and System Dynamics: Ensuring Fairness and Scientific Integrity

The aggregation of individual heroic acts into a cohesive, meaningful Cybernetic Hero score, $H(d)$, for a given identity d, is governed by specific rules designed to ensure systemic fairness, prevent bias, and maintain scientific integrity. The user's clarifications provide two critical constraints: uniform time-weighting for all verified deeds and equal contribution from all duty fields [32] [167]. The principle of uniform time-weighting means that a heroic act performed yesterday is valued identically to one performed ten years ago, provided it has been audited and verified [28]. This deliberate choice removes recency bias, which can plague many reputation systems and favor those active in the present moment regardless of their long-term consistency [144][145]. It incentivizes sustained, lifelong commitment to heroic values rather than short-term, high-profile interventions. While some decay in the perceived relevance of old information is natural, this framework prioritizes historical fidelity and cumulative virtue over immediate impact. This approach is supported by research showing that persistent claims, even without decay, can retain significant weight in complex reasoning models, provided they are consistently validated [28]. The aggregation of these timeless acts can be modeled as a simple sum or average of their respective Output Factors, O_i.

Similarly, the rule of equal contribution from all duty fields—such as PeaceKeeping, EcoCivic, and CivicSupport—ensures that no single domain of heroism is privileged over others [32] [167]. The Social-Impact vector S is designed to capture contributions across these diverse areas, and the final aggregation treats the output of each field equally . This promotes a balanced and holistic view of heroism, recognizing that protecting the environment is as much a heroic contribution as mediating a conflict. This equal weighting is a normative choice aimed at fairness, preventing the system from inadvertently promoting a narrow definition of heroism. Mathematically, this can be achieved by normalizing the scores within each duty field before aggregation or by simply averaging the overall Output Factors, which inherently treats all acts equally regardless of their field of origin [172]. This uniformity simplifies the model and makes it more transparent, reducing the risk of hidden biases in the weighting scheme.

To maintain scientific integrity, the framework incorporates dynamic elements for handling new information. A key feature is the ability to penalize an agent's score when later audits reveal hidden harms in an earlier, previously approved act . This is enforced through a revision process for the associated "EvidenceBundle." If new evidence emerges that contradicts the original claim or reveals unintended negative consequences, the

corresponding act's score can be downgraded or removed. This mirrors practices in digital forensics and legal documentation, where records must be corrected based on new evidence to preserve their integrity [69] [82]. However, this process must be handled with extreme care to avoid accusations of censorship or unfair retroactive punishment. The revision of an EvidenceBundle should be a transparent, auditable event, clearly documented and subject to its own verification process, similar to how revocation lists work for verifiable credentials [2] [4]. The goal is not to erase history but to correct it, ensuring the reputation score remains an accurate and evolving reflection of an identity's conduct. This dynamic updating, combined with uniform weighting and equal field contribution, creates a stable yet responsive aggregation model that is fair, transparent, and grounded in a continuous process of verification and correction.

Operationalization and Practical Challenges: From Theory to Implementation

While the conceptual framework for the Cybernetic Hero score is comprehensive and philosophically robust, its transition from theory to a functional, scalable system faces significant practical challenges related to the operationalization of its metrics, the mitigation of algorithmic bias, and the development of a viable governance process. The most immediate hurdle lies in translating the abstract definitions of the Word-Math and Social-Impact vectors into reliable computational models. Metrics like repetition (`y`) and topic drift (`z`) are relatively straightforward to approximate using techniques from computational linguistics, such as lexical density analysis and semantic similarity measurements [43] [71]. However, more nuanced metrics like kindness (`K`) and the axes of the Social-Impact vector (e.g., `S_antistigma`) are far more complex. The measurement of stigma, for example, suffers from a lack of consensus in the academic literature, with dozens of different definitions and measurement tools, often highly specific to a single disease or population [13]. Developing a generalized, yet sensitive, model for these concepts would require a massive, multidisciplinary effort, integrating insights from psychology, sociology, and computational linguistics. Furthermore, any large-scale system using AI for scoring is susceptible to inheriting and amplifying societal biases present in its training data [39] [144]. A model trained predominantly on data from one cultural context might misinterpret prosocial behavior in another, unfairly penalizing agents from different backgrounds.

Another major challenge is the development of a credible and efficient validation process. The framework's reliance on externally verifiable "EvidenceBundles" is its greatest strength, ensuring scientific validity . However, verifying every act manually is infeasible at scale. A purely automated system risks false positives and negatives, especially with the ethical hard gates. A purely manual system is too slow and prone to human error and bias. The most promising path forward is a hybrid model that combines the strengths of both. Automated systems could be used to pre-screen and flag acts that exhibit characteristics of high toxicity, topic drift, or greed, based on predefined thresholds. These flagged acts would then be routed to a decentralized, peer-reviewed jury of experts or community members for a final decision [45] [108]. This approach leverages machine efficiency for initial triage while retaining human judgment for complex, nuanced, and high-stakes determinations. This model is analogous to approaches used in federated learning, where a global model is updated based on local data, but its trustworthiness depends on a decentralized reputation system to vet the contributions of individual participants [122][127].

Finally, interpretability and usability are critical for the framework's adoption and educational purpose. To understand and train toward a higher score, users need clear, actionable feedback . The system must be designed to explain why a particular narrative received a low K score or a high z score. This "Word-Science" lens involves exposing users to the underlying metrics, allowing them to rewrite their narratives and see the quantitative impact of their revisions . This iterative process turns the framework into a tool for self-improvement, guiding individuals to frame their actions in ways that are not only truthful but also maximize their positive impact and minimize harm. This aligns with findings that multimodal interventions combining psychoeducation with interactive exercises are most effective for changing attitudes [16]. In summary, moving from this compelling vision to a working reality requires a concerted effort in three key areas: (1) developing robust, fair, and auditable computational models for complex metrics; (2) designing a scalable hybrid human-machine governance process for validation; and (3) creating intuitive interfaces that provide meaningful feedback for user training. Addressing these challenges will determine whether the framework remains a theoretical ideal or becomes a practical instrument for fostering a more ethical and impactful society.

Reference

1. Verifiable Credentials Data Model v2.0 - W3C https://www.w3.org/TR/vc-data-model-2.0/
2. CRSet: Private Non-Interactive Verifiable Credential Revocation https://arxiv.org/html/2501.17089v3
3. Verifiable Credentials Data Model 1.0 - W3C https://www.w3.org/TR/2019/CR-verifiable-claims-data-model-20190328/
4. Bitstring Status List v1.0 - W3C https://www.w3.org/TR/vc-bitstring-status-list/
5. (PDF) Personhood credentials: Artificial intelligence and the value of ... https://www.researchgate.net/publication/383153458_Personhood_credentials_Artificial_intelligence_and_the_value_of_privacy-preserving_tools_to_distinguish_who_is_real_online
6. [PDF] Identity in a Digital World A new chapter in the social contract https://www3.weforum.org/docs/WEF_INSIGHT_REPORT_Digital%20Identity.pdf
7. Digital identity: Contemporary challenges for data protection, privacy ... https://www.researchgate.net/publication/333765008_Digital_identity_Contemporary_challenges_for_data_protection_privacy_and_non-discrimination_rights
8. [PDF] Dictionary for Civil Registration and Identification - IADB Publications https://publications.iadb.org/publications/english/document/Dictionary-for-Civil-Registration-and-Identification.pdf
9. Global education monitoring report, 2023: technology in education https://unesdoc.unesco.org/ark:/48223/pf0000385723.
10. Speakers - AI for Good - ITU https://aiforgood.itu.int/summit25/speakers/
11. Civility, Barbarism and the Evolution of International Humanitarian Law https://www.cambridge.org/core/books/civility-barbarism-and-the-evolution-of-international-humanitarian-law/A992B1158546E71DE3704081E53A9E4C
12. [PDF] S/PV.9945 (Resumption 2) Security Council - the United Nations https://docs.un.org/en/S/PV.9945%20(Resumption%202)
13. Stigma measurement in health: a systematic review - PMC https://pmc.ncbi.nlm.nih.gov/articles/PMC12311962/
14. Conceptualizing and Measuring Mental Illness Stigma - PMC - NIH https://pmc.ncbi.nlm.nih.gov/articles/PMC6261312/

15. Anti-stigma interventions in low-income and middle-income countries https://www.thelancet.com/journals/eclinm/article/PIIS2589-5370(24)00191-3/fulltext
16. Evidence for the effectiveness of interventions to reduce ... - BMJ Open https://bmjopen.bmj.com/content/13/2/e067126
17. A Call to Action. A Critical Review of Mental Health Related Anti ... https://www.frontiersin.org/journals/public-health/articles/10.3389/fpubh.2020.569539/full
18. Need or Greed? Conditions for Collective Action against Corruption https://www.researchgate.net/publication/305926983_Need_or_Greed_Conditions_for_Collective_Action_against_Corruption_NEED_OR_GREED
19. The life and times of individuals scoring high and low on ... https://www.sciencedirect.com/science/article/abs/pii/S0092656616300824
20. [PDF] The normalisation of corruption: why it occurs and what can be done ... https://publicadministration.desa.un.org/sites/default/files/old-site/Jon%20Quah%20-%20normalization%20of%20corruption%20-%20final.pdf
21. Does moral commitment predict resistance to corruption ... - PMC - NIH https://pmc.ncbi.nlm.nih.gov/articles/PMC8752004/
22. A Blockchain-Based Foundation for Autonomous AI Agents - arXiv.org https://arxiv.org/html/2602.14219v1
23. DeScAI: the convergence of decentralized science and artificial ... https://www.frontiersin.org/journals/blockchain/articles/10.3389/fbloc.2025.1657050/full
24. [PDF] arXiv:2211.09110v2 [cs.CL] 1 Oct 2023 https://arxiv.org/pdf/2211.09110
25. [PDF] arXiv:2305.10403v3 [cs.CL] 13 Sep 2023 https://arxiv.org/pdf/2305.10403
26. [PDF] Spirituality and Knowledge Dynamics - ResearchGate https://www.researchgate.net/profile/Raysa-Rocha/publication/383276195_Spirituality_and_Knowledge_Dynamics_New_Perspectives_for_Knowledge_Management_and_Knowledge_Strategies/links/66c655be75613475fe70066b/Spirituality-and-Knowledge-Dynamics-New-Perspectives-for-Knowledge-Management-and-Knowledge-Strategies.pdf
27. Psychology Net Notes | PDF - Scribd https://www.scribd.com/document/694553214/Psychology-Net-Notes
28. [PDF] Bayesian Epistemology with Weighted Authority - arXiv.org https://arxiv.org/pdf/2506.16015
29. scrabble.txt - CMU School of Computer Science https://www.cs.cmu.edu/~wlovas/15122-r11/lectures/old/24-tries/scrabble.txt
30. 333333 23135851162 the 13151942776 of 12997637966 ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/words-333333.txt

31. [PDF] WSIS Forum 2019 Outcome Document - ITU https://www.itu.int/en/itu-wsis/Documents/Forum2019/DRAFT-WSISForum2019OutcomeDocument.pdf
32. 3594859 time|3570693 year|3516017 publisher - Apache's svn http://svn.apache.org/repos/asf/lucene/dev/tags/realtime_DWPT_final_2011-05-02/solr/src/test-files/Top50KWiki.utf8
33. 9 - Protecting Warfighters from Superfluous Injury and Unnecessary ... https://www.cambridge.org/core/books/civility-barbarism-and-the-evolution-of-international-humanitarian-law/protecting-warfighters-from-superfluous-injury-and-unnecessary-suffering/36F1FF18D39DD52803491EF1868E7A30
34. (PDF) Compass & Mirror Sexual Markets, Psychoanalysis, and ... https://www.researchgate.net/publication/394940348_Compass_Mirror_Sexual_Markets_Psychoanalysis_and_Cinema_Lacan_Houellebecq_Lynch-and_a_Structural_Re-read_of_Lesbian_Desire
35. Ethical management of human-AI interaction: Theory development ... https://www.sciencedirect.com/science/article/pii/S0963868723000185
36. Ethical Implications and Accountability of Algorithms - ResearchGate https://www.researchgate.net/publication/324896361_Ethical_Implications_and_Accountability_of_Algorithms
37. Ethical Dilemmas and Privacy Issues in Emerging Technologies - PMC https://pmc.ncbi.nlm.nih.gov/articles/PMC9921682/
38. Ethical Decision-Making in Artificial Intelligence: A Logic ... - MDPI https://www.mdpi.com/2673-2688/5/4/130
39. [PDF] The Forgotten Margins of AI Ethics - arXiv https://arxiv.org/pdf/2205.04221
40. Framing and Language of Ethics: Technology, Persuasion, and ... https://www.sciopen.com/article/10.23919/JSC.2021.0027
41. (PDF) THE PRINCIPLES OF DEONTOLOGICAL ETHICS https://www.researchgate.net/publication/381582783_THE_PRINCIPLES_OF_DEONTOLOGICAL_ETHICS
42. [PDF] Data analytics subgroup report - EMA https://www.ema.europa.eu/en/documents/report/data-analytics-subgroup-report_en.pdf
43. Computer Science Nov 2023 - arXiv http://arxiv.org/list/cs/2023-11?skip=1425&show=2000
44. Advanced Sciences and Technologies for Security Applications ... https://www.academia.edu/42727921/Advanced_Sciences_and_Technologies_for_Security_Applications_From_Strategy_to_Implementation
45. Auditing large language models: a three-layered approach https://link.springer.com/article/10.1007/s43681-023-00289-2

46. [PDF] Whitepaper on Ethical Research into Large Language Models - arXiv https://arxiv.org/pdf/2410.19812
47. The 2024 Conference on Empirical Methods in Natural Language ... https://aclanthology.org/events/emnlp-2024/
48. A Comprehensive Survey of Small Language Models in the Era of ... https://dl.acm.org/doi/full/10.1145/3768165
49. Navigating the Ethical Terrain of AI-Generated Text Tools: A Review https://ieeexplore.ieee.org/iel8/6287639/10380310/10813359.pdf
50. [PDF] Global Economic Prospects -- June 2025 - The World Bank https://thedocs.worldbank.org/en/doc/8bf0b62ec6bcb886d97295ad930059e9-0050012025/original/GEP-June-2025.pdf
51. [PDF] The Globalization of Crime - Unodc https://www.unodc.org/roseap/uploads/archive/documents/2010/06/tocta/TOCTA_Report_2010_low_res.pdf
52. Contents of Current Periodicals - jstor https://www.jstor.org/stable/pdf/2564848.pdf
53. [PDF] Medcaptain Medical Technology Co., Ltd. 深圳麥科田生物醫療技術 ... https://www1.hkexnews.hk/app/sehk/2025/107676/documents/sehk25091100252.pdf
54. (PDF) Mapping the Ethicality of Algorithmic Pricing: A Review of ... https://www.researchgate.net/publication/337879365_Mapping_the_Ethicality_of_Algorithmic_Pricing_A_Review_of_Dynamic_and_Personalized_Pricing
55. a modular framework for ethical, structured, and adaptive AI - Springer https://link.springer.com/article/10.1007/s10462-025-11330-7
56. Algorithmic discrimination in the credit domain: what do we know ... https://www.researchgate.net/publication/370839603_Algorithmic_discrimination_in_the_credit_domain_what_do_we_know_about_it
57. [PDF] The impact of Artificial Intelligence on communication networks ... - ITU https://www.itu.int/dms_pub/itu-s/opb/journal/S-JOURNAL-ICTF.VOL1-2018-1-PDF-E.pdf
58. Shaping capabilities with token-level data filtering - arXiv https://arxiv.org/html/2601.21571v1
59. Shaping capabilities with token-level data filtering - ResearchGate https://www.researchgate.net/publication/400237731_Shaping_capabilities_with_token-level_data_filtering
60. [PDF] Mind as Theory Engine: Causation, Explanation and Time https://escholarship.org/content/qt4nk6p0f3/qt4nk6p0f3.pdf

61. English Text (1.77 MB) - World Bank Open Knowledge Repository https://openknowledge.worldbank.org/bitstreams/8615f608-a7be-54e7-9c47-70ee45686791/download
62. Statistical Approaches (Part IV) - The Cambridge Handbook of ... https://www.cambridge.org/core/books/cambridge-handbook-of-research-methods-and-statistics-for-the-social-and-behavioral-sciences/statistical-approaches/0E2B991A5022167B063091E16333E92B
63. 424B4 - SEC.gov https://www.sec.gov/Archives/edgar/data/1997464/000119312525083415/d943351d424b4.htm
64. [PDF] Civil and military responsibilities in disrupted states https://digitallibrary.un.org/record/504689/files/Fromcivilstrife.pdf
65. [PDF] The Globalization of Crime - Unodc https://www.unodc.org/documents/lpo-brazil/noticias/2010/06/TOCTA_Report_2010_low_res.pdf
66. Digital Targeting: Artificial Intelligence, Data, and Military Intelligence https://academic.oup.com/jogss/article/9/2/ogae009/7667104
67. Evidentiality and epistemic modality: On the close relationship ... https://www.researchgate.net/publication/233521897_Evidentiality_and_epistemic_modality_On_the_close_relationship_between_two_different_categories
68. Organizational Justice and Behavioral Ethics: Promises and Prospects https://www.researchgate.net/publication/261251235_Organizational_Justice_and_Behavioral_Ethics_Promises_and_Prospects
69. [PDF] Fraud Ethics Integrity and and Corruption - European Court of Auditors https://www.eca.europa.eu/lists/ecadocuments/journal19_02/journal19_02.pdf
70. Proceedings of the 31st International Conference on Computational ... https://aclanthology.org/volumes/2025.coling-main/
71. Proceedings of the 61st Annual Meeting of the Association for ... https://aclanthology.org/volumes/2023.acl-long/
72. Hw3 Stats Google 1gram | PDF | Internet Forum | Software - Scribd https://www.scribd.com/document/917428277/Hw3-Stats-Google-1gram
73. International Conference on Computational Linguistics (2022) https://aclanthology.org/events/coling-2022/
74. [PDF] EU Drug Markets Report 2019 https://www.euda.europa.eu/system/files/publications/12078/20192630_TD0319332ENN_PDF.pdf
75. Implementing Climate Change Policy - Springer Link https://link.springer.com/content/pdf/10.1007/978-981-95-0607-1.pdf
76. The 31st International Conference on Computational Linguistics https://aclanthology.org/events/coling-2025/

77. 自然语言处理2025_5_21 - arXiv每日学术速递 https://www.arxivdaily.com/thread/67649
78. The Falcon Series of Open Language Models - arXiv https://arxiv.org/html/2311.16867
79. [PDF] GPT-4 Technical Report - AMiner https://static.aminer.cn/upload/pdf/1792/141/137/641130e378d68457a4a2986f_0.pdf
80. A large-scale audit of dataset licensing and attribution in AI https://www.researchgate.net/publication/383566266_A_large-scale_audit_of_dataset_licensing_and_attribution_in_AI
81. [PDF] Open Problems in Technical AI Governance - OpenReview https://openreview.net/pdf?id=1nO4qFMiS0
82. [PDF] Conference Proceedings >> Download - ijsrst https://ijsrst.com/paper/v12i9.pdf
83. [PDF] Handbook of Theories of Social Psychology, Volume 1 http://www.personpsy.org/uploadfiles/file/books/Handbook%20of%20Theories%20of%20Social%20Psychology%20Volume%20One%20(Paul%20A.%20M.%20Van%20Lange,%20Arie%20W.%20Kruglanski%20etc.)%20(Z-Library)(1).pdf
84. (PDF) AI AND EMOTIONS IN DIGITAL SOCIETY - Academia.edu https://www.academia.edu/122960609/AI_AND_EMOTIONS_IN_DIGITAL_SOCIETY
85. Download book PDF - Springer https://link.springer.com/content/pdf/10.1007/978-1-4612-3628-3.pdf
86. [PDF] The Sandflow Construction Theory: A Comprehensive Engineering ... https://zenodo.org/records/18036910/files/The%20Sandflow%20Construction%20Theory.pdf?download=1
87. Verifiable Credentials Data Model v2.0 - W3C https://www.w3.org/TR/2023/WD-vc-data-model-2.0-20230711/
88. (PDF) Customer-Centric AI: Conversational Technologies ... https://www.researchgate.net/publication/396160996_Customer-Centric_AI_Conversational_Technologies_Personalization_and_Ethical_Innovation
89. Entries on Transitional Justice Debates, Controversies, and Key ... https://www.cambridge.org/core/books/encyclopedia-of-transitional-justice/entries-on-transitional-justice-debates-controversies-and-key-questions/678FB58E1B509A74A0E9118B140F1B67
90. Lawyers' Ethics and Professional Regulation 4th Edition | PDF - Scribd https://www.scribd.com/document/770682126/Lawyers-Ethics-and-Professional-Regulation-4th-Edition

91. Evidentiality and Its Relation to Epistemic Modality - ResearchGate https://www.researchgate.net/publication/239921261_Semantic_Maps_and_the_Identification_of_Cross-Linguistic_Generic_Categories_Evidentiality_and_Its_Relation_to_Epistemic_Modality
92. (PDF) The Character of Epistemic Modality: Evidential indexicals https://www.academia.edu/42728051/The_Character_of_Epistemic_Modality_Evidential_indexicals
93. Epistemicity, Evidentiality, and Theory of Mind (ToM) - ResearchGate https://www.researchgate.net/publication/333031693_Epistemicity_Evidentiality_and_Theory_of_Mind_ToM
94. Methodological Problems in The Analysis | PDF | Linguistics - Scribd https://www.scribd.com/document/384983949/Methodological-Problems-in-the-Analysis-1
95. [PDF] DeepSeek-R1 Thoughtology: Let's think about LLM reasoning https://openreview.net/pdf/807fd0c52fde8e65b870d177c0d5992d3dc08552.pdf
96. (PDF) The bias beneath: analyzing drift in YouTube's algorithmic ... https://www.researchgate.net/publication/383395529_The_bias_beneath_analyzing_drift_in_YouTube's_algorithmic_recommendations
97. Why is constrained neural language generation particularly ... - arXiv https://arxiv.org/html/2206.05395v2
98. [PDF] Beyond Refusals: AQI as an Intrinsic Alignment - ACL Anthology https://aclanthology.org/2025.emnlp-main.145.pdf
99. [PDF] Overview Methodology Frameworks https://13115299.s21i.faiusr.com/61/1/ABUIABA9GAAgk9j8jQYo-KyP7gM.pdf
100. 插件8：拼写检查 - CSDN博客 https://blog.csdn.net/manbujingxin/article/details/6845477
101. Abstractive - Nicholas Carlini https://nicholas.carlini.com/writing/2019/advex_papers_with_abstracts.txt
102. Findings of the Association for Computational Linguistics: ACL 2023 https://aclanthology.org/volumes/2023.findings-acl/
103. Findings of the Association for Computational Linguistics (2023) https://aclanthology.org/events/findings-2023/
104. data/wordlists/password.lst - GitLab https://git.ustc.edu.cn/cwzsquare/metasploit-framework/-/blob/60210f57e97119ed7faad1ba6a08720d64601747/data/wordlists/password.lst
105. (PDF) English Homophones and Spelling - Academia.edu https://www.academia.edu/36558860/English_Homophones_and_Spelling

106. Verifiable Credentials Data Model v2.0 - W3C https://www.w3.org/TR/2022/WD-vc-data-model-2.0-20220811/

107. Deontic Temporal Logic for Formal Verification of AI Ethics https://www.researchgate.net/publication/387954314_Deontic_Temporal_Logic_for_Formal_Verification_of_AI_Ethics

108. A Formal Verification of a Reputation Multi-Factor Authentication ... https://www.mdpi.com/1424-8220/23/15/6933

109. [PDF] Understanding algorithmic decision-making: Opportunities and ... https://www.europarl.europa.eu/RegData/etudes/STUD/2019/624261/EPRS_STU(2019)624261_EN.pdf

110. Clarifying Assumptions About Artificial Intelligence Before ... https://academic.oup.com/grurint/article/71/4/295/6528412

111. Part I - AI, Ethics and Philosophy https://www.cambridge.org/core/books/cambridge-handbook-of-the-law-ethics-and-policy-of-artificial-intelligence/ai-ethics-and-philosophy/EFEFAE5CD989984FD6C3FA64BFB25658

112. Download book PDF - Springer Link https://link.springer.com/content/pdf/10.1057/9781137281555.pdf

113. 插件8：拼写检查转载 - CSDN博客 https://blog.csdn.net/dearbaba_1666/article/details/80829311

114. Arxiv今日论文| 2025-12-23 - 闲记算法 http://lonepatient.top/2025/12/23/arxiv_papers_2025-12-23

115. Hw3-stats-google-1gram - CMU School of Computer Science https://www.cs.cmu.edu/～roni/11761/2017_fall_assignments/hw3_stats_google_1gram.txt

116. Evidentiality: Linguistic categories and grammaticalization https://www.researchgate.net/publication/233611986_Evidentiality_Linguistic_categories_and_grammaticalization

117. [PDF] Evidentiality, egophoricity, and engagement - Zenodo https://zenodo.org/record/3968344/files/261.pdf?download=1

118. (PDF) Routledge Handbook of Surveillance Studies - Academia.edu https://www.academia.edu/6778392/Routledge_Handbook_of_Surveillance_Studies

119. [PDF] The Politics of Representation and the Social Order - eScholarship.org https://escholarship.org/content/qt4xt6g5v1/qt4xt6g5v1_noSplash_af023b9531b9ff74541e3b5bec01914f.pdf

120. Investigative Interviewing - Springer https://link.springer.com/content/pdf/10.1007/978-1-4614-9642-7.pdf

121. RTIFed: A Reputation-Based Triple-Step Incentive Mechanism for ... https://arxiv.org/html/2505.13319v1

122. [PDF] Trustworthy Federated Learning: A Survey - arXiv https://arxiv.org/pdf/2305.11537
123. Zero-Knowledge Proof-based Verifiable Decentralized Machine ... https://arxiv.org/html/2310.14848v2
124. Modular Federated Learning: A Meta-Framework Perspective - arXiv https://arxiv.org/html/2505.08646v1
125. Mixture of Experts for Decentralized Generative AI and ... - arXiv.org https://arxiv.org/html/2504.19660v2
126. [PDF] privacy-preserving in blockchain-based federated learning systems https://arxiv.org/pdf/2401.03552
127. [PDF] Federated Learning Meets Blockchain in Edge Computing - arXiv.org https://arxiv.org/pdf/2104.01776
128. Advances and Open Challenges in Federated Learning with ... - arXiv https://arxiv.org/html/2404.15381v2
129. Federated Learning for Cyber Physical Systems: A Comprehensive ... https://arxiv.org/html/2505.04873v1
130. 1 Introduction - arXiv https://arxiv.org/html/2602.05289v1
131. [PDF] 2026 POMS-HK Book of Abstracts https://sme.cuhk.edu.cn/sites/default/files/2025-12/Book_of_Abstracts_12_25_v1.pdf
132. (PDF) Transaction Fee Mechanisms and User Welfare in ... https://www.researchgate.net/publication/393628824_Transaction_Fee_Mechanisms_and_User_Welfare_in_Permissionless_Blockchains
133. Challenges to the impartiality and identity of humanitarian action https://www.cambridge.org/core/journals/international-review-of-the-red-cross/article/imperfect-relief-challenges-to-the-impartiality-and-identity-of-humanitarian-action/0520C35DD2132B3135BFC635098C6F7B
134. Threats and Defenses in Federated Learning Life Cycle - arXiv https://arxiv.org/html/2407.06754v1
135. Gaming and Cooperation in Federated Learning: What Can Happen ... https://arxiv.org/html/2509.02391
136. [PDF] Threats and Defenses in Federated Learning Life Cycle - arXiv https://arxiv.org/pdf/2407.06754
137. [PDF] FMore: An Incentive Scheme of Multi-dimensional Auction for ... - arXiv https://arxiv.org/pdf/2002.09699
138. A Survey of Trustworthy Federated Learning: Issues, Solutions, and ... https://dl.acm.org/doi/full/10.1145/3678181

139. [PDF] Synergies between machine learning and reasoning https://www.sciencedirect.com/science/article/am/pii/S0888613X24000938
140. Tactics for Internal Compliance: A Literature Review - ResearchGate https://www.researchgate.net/publication/343568542_Tactics_for_Internal_Compliance_A_Literature_Review
141. [PDF] Synergies between machine learning and reasoning - HAL https://hal.science/hal-04653616/file/1-s2.0-S0888613X24000938-main.pdf
142. [PDF] From Rules to Language Models - ACL Anthology https://aclanthology.org/2025.r2lm-1.pdf
143. A Systematic Review of Blockchain-Based Privacy-Preserving ... https://dl.acm.org/doi/full/10.1145/3674156
144. Socioeconomic Threats of Deepfakes and the Role of Cyber ... https://mags.acm.org/communications/september_2025/MobilePagedArticle.action?articleId=2079794
145. An Analysis of the Impact of Gold Open Access Publications in ... https://mags.acm.org/communications/september_2025/MobilePagedArticle.action?articleId=2079793
146. Reinforcement Learning: An Introduction | Guide books https://dl.acm.org/doi/book/10.5555/3312046
147. Anomaly Detection in Restaurant Receipts Data - IEEE Xplore https://ieeexplore.ieee.org/iel8/6287639/10380310/10701309.pdf
148. Computers and Intractability: A Guide to the Theory of NP ... - ACM https://dl.acm.org/doi/10.5555/578533
149. Fairness-Driven Recommendation in Heterogeneous Two-Sided ... https://arxiv.org/html/2602.10739v1
150. Dynamic Fairness-aware Recommendation Through Multi-agent ... https://arxiv.org/html/2303.00968v3
151. [PDF] Fair Classification with Noisy Protected Attributes - arXiv.org https://arxiv.org/pdf/2006.04778
152. Fair Multi-agent Persuasion with Submodular Constraints - arXiv https://arxiv.org/html/2511.08538v1
153. Multi-dimensional Discrimination in Law and Machine Learning https://dl.acm.org/doi/fullHtml/10.1145/3593013.3593979
154. Dynamic Fairness-aware Recommendation Through Multi-agent ... https://dl.acm.org/doi/full/10.1145/3690653
155. A Theoretical Framework for Adaptive Utility-Weighted Benchmarking https://arxiv.org/html/2602.12356v1

156. [PDF] Fair Representation: Guaranteeing Approximate Multiple Group ... https://arxiv.org/pdf/2109.00545

157. [PDF] Understanding and Improving Fairness-Accuracy Trade-offs in Multi ... https://arxiv.org/pdf/2106.02705

158. MAFE: Enabling Equitable Algorithm Design in Multi-Agent ... - arXiv https://arxiv.org/html/2502.18534v2

159. WWW '23: Proceedings of the ACM Web Conference 2023 - SIGWEB https://www.sigweb.sigweb.hosting.acm.org/toc/www23.html

160. Vol. 14 :: No. 1 :: Jan – Mar 2019 - IEEE IC Vaani https://r10.ieee.org/indiacouncil/wp-content/uploads/sites/149/2019/04/icnl-19q1.pdf

161. [PDF] Dynamic Fairness-aware Recommendation Through Multi-agent ... https://arxiv.org/pdf/2303.00968

162. [PDF] arXiv:2411.02654v1 [cs.GT] 4 Nov 2024 https://arxiv.org/pdf/2411.02654?

163. Ethics, Governance, and Fairness in Large-Scale Data Science and ... https://www.researchgate.net/publication/397038403_Ethics_Governance_and_Fairness_in_Large-Scale_Data_Science_and_Algorithmic_Decision_Systems

164. [PDF] Assessing-the-Applicability-of-Participatory-Multi-Criteria-Analysis ... https://www.researchgate.net/profile/Marco-Dean/publication/343722759_Assessing_the_Applicability_of_Participatory_Multi-Criteria_Analysis_Methodologies_to_the_Appraisal_of_Mega_Transport_Infrastructure/links/5f3bf8a492851cd302019586/Assessing-the-Applicability-of-Participatory-Multi-Criteria-Analysis-Methodologies-to-the-Appraisal-of-Mega-Transport-Infrastructure.pdf

165. [PDF] GAUSS: Benchmarking Structured Mathematical Skills for Large ... https://arxiv.org/pdf/2509.18122

166. (PDF) Multi-Criteria Analysis: A Manual - ResearchGate https://www.researchgate.net/publication/30529921_Multi-Criteria_Analysis_A_Manual

167. Optimizing data-driven weights in multidimensional indexes https://www.researchgate.net/publication/394602257_Optimizing_data-driven_weights_in_multidimensional_indexes

168. A decision architecture for epistemic prioritization: Machine learning ... https://www.sciencedirect.com/science/article/pii/S0160791X25002295

169. [PDF] Ranking with Fairness Constraints - arXiv https://arxiv.org/pdf/1704.06840

170. (PDF) Analyzing Incentives and Fairness in Ordered Weighted ... https://arxiv.org/pdf/2410.12884

171. Why are there many equally good models? An Anatomy of ... - arXiv https://arxiv.org/html/2601.06730v2

172. Supplementary Material for AgentScore - arXiv.org https://arxiv.org/html/2601.22324v1
173. Towards the Formalization of a Trustworthy AI for Mining ... https://arxiv.org/html/2510.20621v2
174. Multi-party Computation Protocols for Post-Market Fairness ... https://arxiv.org/html/2602.01837v1
175. Fair Document Valuation in LLM Summaries via Shapley Values https://arxiv.org/html/2505.23842v3
176. [PDF] Multi-party Computation Protocols for Post-Market Fairness ... https://www.arxiv.org/pdf/2602.01837
177. ANML: Attribution-Native Machine Learning with Guaranteed ... - arXiv https://arxiv.org/html/2602.11690v1
178. Value Alignment of Social Media Ranking Algorithms - arXiv https://arxiv.org/html/2509.14434v1
179. [PDF] Dynamic Delegation with Reputation Feedback - arXiv https://arxiv.org/pdf/2508.19676
180. Dynamic Delegation with Reputation Feedback † - † thanks - arXiv.org https://arxiv.org/html/2508.19676v2
181. Smart Network Field Theory: The Technophysics of Blockchain and ... https://arxiv.org/pdf/1810.09514
182. [PDF] Some Simple Economics of AGI - arXiv https://www.arxiv.org/pdf/2602.20946
183. Some Simple Economics of AGI † - † thanks - arXiv https://arxiv.org/html/2602.20946v1
184. Artificial Intelligence - arXiv.org https://arxiv.org/list/cs.AI/new
185. Computer Science Apr 2024 - arXiv https://www.arxiv.org/list/cs/2024-04?skip=6000&show=2000
186. Artificial Intelligence Jun 2025 - arXiv.org https://www.arxiv.org/list/cs.AI/2025-06?skip=2900&show=2000
187. Computer Science Aug 2020 - arXiv http://arxiv.org/list/cs/2020-08?skip=3475&show=2000
188. Computer Science Jan 2022 - arXiv http://arxiv.org/list/cs/2022-01?skip=500&show=2000
189. Computer Science 2016 - arXiv https://www.arxiv.org/list/cs/2016?skip=3100&show=2000
190. Computer Science Jun 2020 - arXiv http://arxiv.org/list/cs/2020-06?skip=4465&show=2000

191. Arxiv今日论文| 2026-02-16 | 闲记算法 http://lonepatient.top/2026/02/16/arxiv_papers_2026-02-16

192. Fifty years of multi-objective optimization and decision-making https://www.sciencedirect.com/science/article/pii/S0377221725004849

193. Efficient Attention Mechanisms for Large Language Models: A Survey https://arxiv.org/html/2507.19595v3

194. a robust evidential methodology for dimensional accuracy and ... https://link.springer.com/article/10.1007/s00170-026-17524-9

195. Robust Disturbance-Response Feature Modeling and Multi ... - MDPI https://www.mdpi.com/2227-7390/14/2/316

196. A Robust Hybrid Machine Learning Framework for Multidimensional ... https://www.sciencedirect.com/science/article/pii/S2590123026005414

197. (PDF) A Comparative Analysis of Multi-Criteria Decision-Making ... https://www.researchgate.net/publication/388553191_A_Comparative_Analysis_of_Multi-Criteria_Decision-Making_Methods_and_Normalization_Techniques_in_Holistic_Sustainability_Assessment_for_Engineering_Applications

198. Some Statistical Approaches to Multidimensional Scaling Data - jstor https://www.jstor.org/stable/2981865